

Section V: Random Forests and GAMs

450C

Stanford University

Department of Political Science

Toby Nowacki

Zuhad Hai

Overview

1. Overview
2. Random Forests: Intuition
3. Random Forests: Implementation
4. GAMs: Intuition
5. GAMs: Implementation
6. Midterm Revision

Random Forests: Intuition

- We want to compute the average \bar{y} for every partition of the data, where the partition is a unique combination of covariates.
- Why is the curse of dimensionality a problem here?

Random Forests: Intuition

- Random Forests give us a way out by searching for the best way to split the multidimensional space
- Within each region, compute the average value of y
- But how to find optimal region?
- Greedy algorithm: tries to find partition that satisfies local minimum of prediction error
- What can go wrong with the greedy algorithm?

Random Forests: Intuition

- To mitigate concern, we introduce random sampling across variables (select z of the J variables)
- When different variables are selected, we will also observe different nodes / trees!
- In general, no good advice on how deep we should grow these trees / how many trees we want
- Tree depth comes at a bias-variance tradeoff: the less data we have in each node, the more do we run the risk of overfitting.
- Can do crossvalidation!

Random Forests: Implementation

Let's prepare our data.

```
library(randomForest)
library(mlbench)
library(caret)

data(Sonar)
df ← Sonar
x ← df[, 1:50]
y ← df[, 51]
```

Random Forests: Implementation

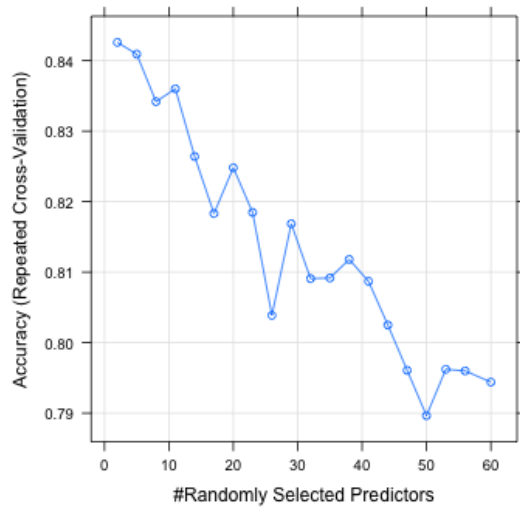
Fit the model.

```
set.seed(2020)
control ← trainControl(method = "repeatedcv",
                        number = 10, repeats = 3)
metric ← "Accuracy"
rf_random ← train(Class ~ ., data = df, method = "rf",
                  metric = metric, tuneLength = 20, trControl = control)
```


Random Forests: Implementation

Accuracy by tree length:

```
plot(rf_random)
```



Random Forests: Implementation

```
tg ← expand.grid(.mtry = c(10:20))  
rf_grid ← train(Class ~ ., data = df, method = "rf",  
  metric = metric, tuneGrid = tg, trControl = control)
```

Random Forests: Implementation

```
print(rf_grid)
```

```
## Random Forest
##
## 208 samples
## 60 predictor
## 2 classes: 'M', 'R'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 188, 188, 187, 187, 187, 187, ...
## Resampling results across tuning parameters:
##
##  mtry  Accuracy   Kappa
##  10    0.8366089  0.6685702
##  11    0.8368543  0.6695420
##  12    0.8415296  0.6787108
##  13    0.8286003  0.6521362
##  14    0.8287518  0.6527369
##  15    0.8223160  0.6390713
##  16    0.8207359  0.6358603
##  17    0.8173232  0.6290108
##  18    0.8140765  0.6217789
##  19    0.8124820  0.6193741
```

Generalised Additive Models: Intuition

- GAMs introduce non-linearity into our classic regression framework:

$$y_i = \beta_0 + s_1(x_{1i}) + s_2(x_{2i}) + s_3(x_{3i}) + u_i$$

where the functions s_1 etc. are estimated from the data.

- Theory somewhat involved, but the key takeaway is that we rely on partial residuals (the relationship between x_1 and y after controlling for the rest)
- GAMs allow us to interpret the relationship between any variable and the outcome in a bivariate plot
- Crucial to remember that the plots show changes in y *relative to its mean*
- Interactions can be modelled with GAMs, but quickly run into the curse of dimensionality problem again.

Generalised Additive Models: Implementation

- Let's compare OLS and GAM results.
- Data and example taken from Peisakhin and Rozenas (2018)
- How does exposure to Russian propaganda media sources affect political behaviour?

```
d ← read.csv("data.csv")
d ← na.omit(d)

head(d)
```

```
## precinct oblast places noplaces type size ukrainian district14par
## 1 590884 Сумська СУМИ 1 city 3 77.44 157
## 2 590885 Сумська СУМИ 1 city 3 77.44 157
## 3 590886 Сумська СУМИ 1 city 3 77.44 157
## 4 590887 Сумська СУМИ 1 city 3 77.44 157
## 5 590888 Сумська СУМИ 1 city 3 77.44 157
## 6 590889 Сумська СУМИ 1 city 3 77.44 157
## registered14par voted14parl oppblock14par porosh14par r14parl turnout14
## 1 1552 844 0.04976303 0.2500000 13.98104 0.543
## 2 2368 1370 0.04087591 0.2503650 11.24088 0.578
## 3 1564 887 0.03720406 0.2559188 11.49944 0.567
## 4 2152 1252 0.05191693 0.2739617 11.50160 0.581
```

Generalised Additive Models: Implementation

```
form0 <- formula("r14pres ~ qualityq + distrussia + factor(Raion) + u")
m0 <- lm(form0, data = d)
coeftest(m0, vcovCL(m0, cluster = m0$model[["factor(Raion)"]]))
```


```
##
## t test of coefficients:
##
##
```

	Estimate	Std. Error	t value	Pr(> t)	
## (Intercept)	4.619296	5.363981	0.8612	0.3892039	
## qualityq	6.431126	2.616363	2.4580	0.0140180	*
## distrussia	-2.012976	0.725696	-2.7739	0.0055690	**
## factor(Raion)Balakliis	20.586593	3.044345	6.7622	1.586e-11	***
## factor(Raion)Barvinkivs	13.964029	2.717876	5.1378	2.931e-07	***
## factor(Raion)Bilopils	-7.047792	1.514835	-4.6525	3.401e-06	***
## factor(Raion)Blyzniukivs	7.403772	2.365735	3.1296	0.0017650	**
## factor(Raion)Bobrovyts	-0.717404	0.311357	-2.3041	0.0212743	*
## factor(Raion)Bohoduhiivs	8.789799	2.669802	3.2923	0.0010036	**
## factor(Raion)Borivs	8.143906	2.761783	2.9488	0.0032114	**
## factor(Raion)Borznians	-0.310447	0.328731	-0.9444	0.3450404	
## factor(Raion)Buryns	-4.047601	0.919786	-4.4006	1.112e-05	***
## factor(Raion)Chernihivs	-2.687138	0.646097	-4.1590	3.273e-05	***
## factor(Raion)Chuhuiv	20.010818	2.723104	7.3485	2.483e-13	***
## factor(Raion)Dorohobuzh	16.402021	2.501010	6.5585	6.227e-11	***

##

Generalised Additive Models: Implementation

```
form1 <- formula("r14pres ~ qualityq + s(distrussia) + factor(Raion)  
m1 <- gam(form1, data = d)  
plot(m1)
```



Generalised Additive Models: Implementation

```
form2 ← formula("r14pres ~ qualityq + s(distrussia) + s(ukrainian) +  
m2 ← gam(form2, data = d)  
plot(m2)
```


Midterm revision

- What have we covered so far?
 - Maximum Likelihood
 - Probit and Logit: Estimation and Uncertainty
 - Principal Components Analysis
 - Ridge, LASSO and Naive Bayes
 - Random Forests, Ensemble Methods and GAMs
- You should be comfortable with:
 - fitting these models to data
 - interpreting the model output
 - evaluating the model's fit, strengths and weaknesses
 - critically applying these techniques to new problems
- We do **not** expect you to:
 - solve complex algebra or other mathematical problems
 - develop new code for applications outside of class